

# An Introduction to Deep Learning

(part 3)




Farnoosh Heidarivincheh  
EMAT31530 - February 2021

# Convolutional Neural Networks

So far, we have discussed

- Neural networks with Fully-Connected layers
- The optimization process / How the network is trained



# Outline

- Convolutional Neural Networks
  - Different Layers
  - Popular Architectures
- Other Popular Neural Networks
  - Recurrent Neural Networks
  - Long Short-Term Memory Networks
  - Auto-Encoders
  - Generative Adversarial Networks

# Convolutional Neural Networks


# Convolutional Neural Networks

## Image classification with **CNN**

- No need to pre-process images
- CNNs outperform traditional approaches





Convolutional  
Neural Network  
(CNN)



# Convolutional Neural Networks

**Image** classification with **CNN** - Layers in a CNN:


- Input
- Convolutional
- Activation (ReLU, ...)
- Pooling
- Fully-Connected (FC)
- Drop-out
- Batch Normalisation (BN), ...



# Convolutional Neural Networks

## Convolutional layer


- It is the main building block of a CNN
- its neurons can have more than one dimension
- It captures the correlation among neighboring data points




# Convolutional Neural Networks

## Convolutional layer

RGB image input  $\mathbf{X}$



Convolution kernel  $\mathbf{W}$



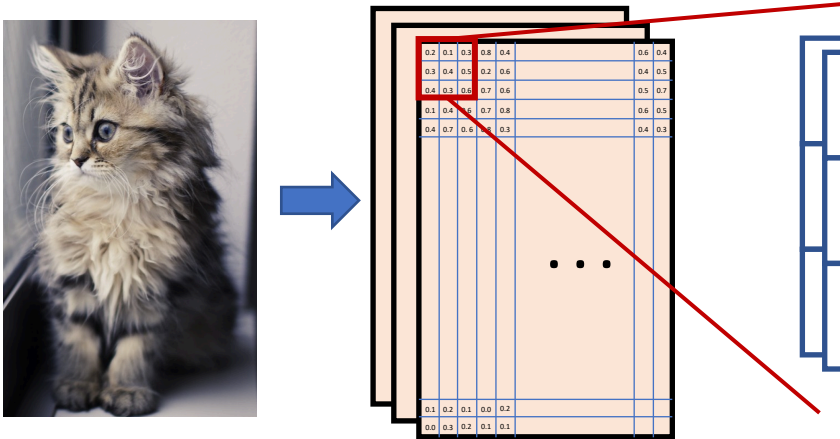
Kernel size:  $3 \times 3 \times 3$



# Convolutional Neural Networks

## Convolutional layer

RGB image input  $\mathbf{X}$



0.2	0.1	0.3
0.3	0.4	0.5
0.4	0.3	0.6

Convolution kernel  
 $\mathbf{W}$


$w_{-1,-1,-1}$	$w_{-1,0,-1}$	$w_{-1,1,-1}$
$w_{0,-1,-1}$	$w_{0,0,-1}$	$w_{0,1,-1}$
$w_{1,-1,-1}$	$w_{1,0,-1}$	$w_{1,1,-1}$

Kernel size:  $3 \times 3 \times 3$

# Convolutional Neural Networks


## Convolutional layer

RGB image input  $\mathbf{X}$



0.2	0.1	0.3
0.3	0.4	0.5
0.4	0.3	0.6


$$w_b + \sum_{l,m,n=-1}^1 x_{i+l,j+m,k+n} w_{l,m,n}$$



# Convolutional Neural Networks

## Convolutional layer


RGB image input  $\mathbf{X}$



0.2	0.1	0.3
0.3	0.4	0.5
0.4	0.3	0.6

$$w_b + \sum_{l,m,n=-1}^1 x_{i+l,j+m,k+n} \underline{w_{l,m,n}}$$


$\mathbf{w}$



# Convolutional Neural Networks

## Convolutional layer


RGB image input  $\mathbf{X}$



0.2	0.1	0.3
0.3	0.4	0.5
0.4	0.3	0.6

$$w_b + \sum_{l,m,n=-1}^1 \underline{x_{i+l,j+m,k+n}} w_{l,m,n}$$


$\mathbf{w}$



# Convolutional Neural Networks

## Convolutional layer

RGB image input  $\mathbf{X}$



0.2	0.1	0.3
0.3	0.4	0.5
0.4	0.3	0.6


$n$	$m$	
$l$		
		$w_{-1,-1,-1}$
		$w_{-1,0,-1}$
		$w_{-1,1,-1}$
		$w_{0,-1,-1}$
		$w_{0,0,-1}$
		$w_{0,1,-1}$
		$w_{1,-1,-1}$
		$w_{1,0,-1}$
		$w_{1,1,-1}$

$$\text{bias } \boxed{w_b} + \sum_{l,m,n=-1}^1 x_{i+l,j+m,k+n} w_{l,m,n}$$

# Convolutional Neural Networks

## Convolutional layer


RGB image input  $\mathbf{X}$



0.2	0.1	0.3
0.1	0.4	0.5
0.4	0.3	0.6

$$w_b + 0.2 \times w_{-1,-1,-1}$$


$$w_b + \sum_{l,m,n=-1}^1 x_{i+l,j+m,k+n} w_{l,m,n}$$



# Convolutional Neural Networks

## Convolutional layer

RGB image input  $\mathbf{X}$



$$w_b + \sum_{l,m,n=-1}^1 x_{i+l,j+m,k+n} w_{l,m,n}$$

0.2	0.1	0.3
0.1	0.4	0.5
0.4	0.3	0.6

$\mathbf{W}$

Diagram illustrating a 3D weight tensor  $\mathbf{W}$  with dimensions  $n$  (height),  $m$  (width), and  $l$  (depth). The tensor is represented as a stack of  $l$   $n \times m$  matrices. The first matrix is shown with weights  $w_{i,j,k}$  for  $i \in \{-1, 0, 1\}$  and  $j \in \{-1, 0, 1\}$ .

$w_{-1,-1,-1}$	$w_{-1,0,-1}$	$w_{-1,1,-1}$
$w_{0,-1,-1}$	$w_{0,0,-1}$	$w_{0,1,-1}$
$w_{1,-1,-1}$	$w_{1,0,-1}$	$w_{1,1,-1}$

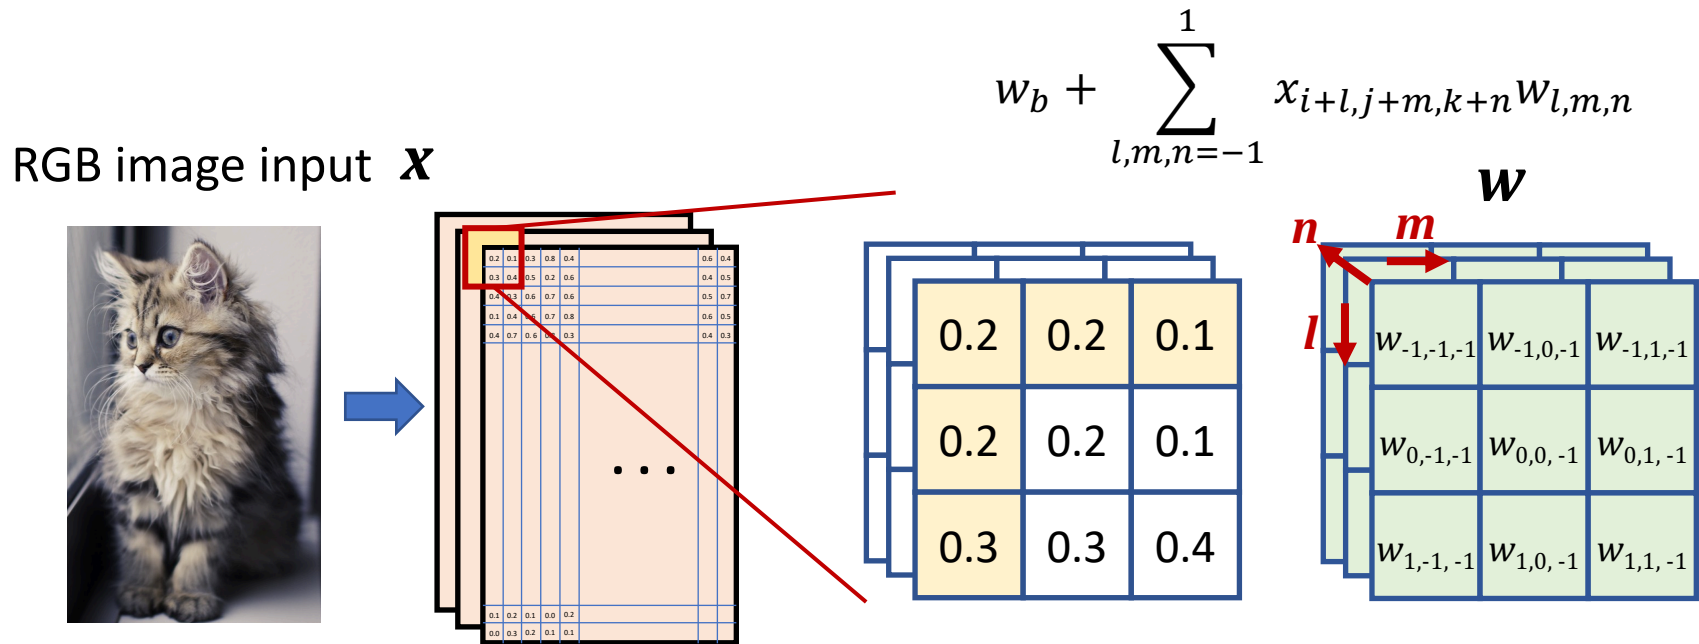
This is repeated for every pixel of the image (if *stride*=1)

Image is padded for the bordering pixels

So that the output is the same size as its input

# Convolutional Neural Networks

## Convolutional layer



This is repeated for every pixel of the image (if *stride*=1)

Image is padded for the bordering pixels

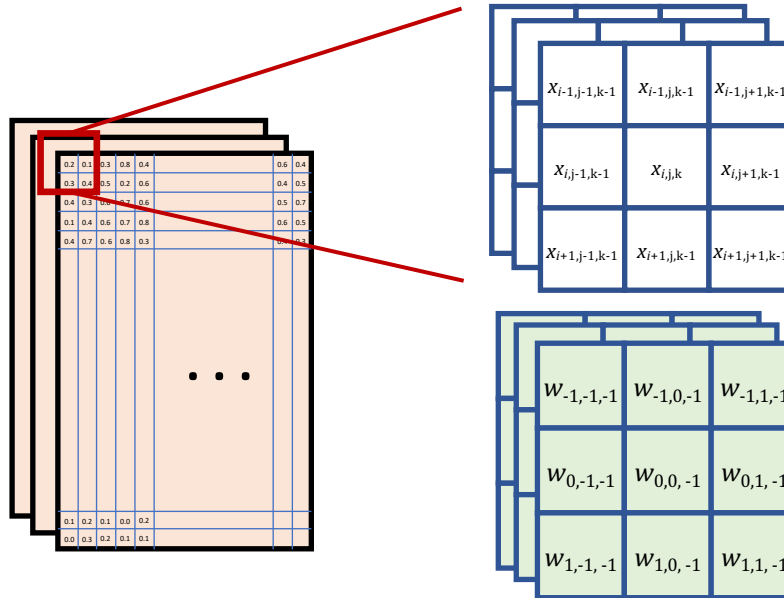
So that the output is the same size as its input




# Convolutional Neural Networks

## Convolutional layer

RGB image input



Conv layer output



This is repeated for every pixel of the image (if *stride*=1)



Image is padded for the bordering pixels

So that the output is the same size as its input


# Convolutional Neural Networks

## Convolutional layer

RGB image input



Conv layer output



This is repeated for every pixel of the image (if *stride*=1)


Image is padded for the bordering pixels

So that the output is the same size as its input

# Convolutional Neural Networks

## Convolutional layer

RGB image input  $\mathbf{x}$




$\mathbf{w}^{(1)}$



$\mathbf{w}^{(2)}$



$\mathbf{w}^{(3)}$




⋮


$\mathbf{w}^{(c-1)}$



$\mathbf{w}^{(c)}$




Conv layer output




# Convolutional Neural Networks

## Convolutional layer

RGB image input  $\mathbf{x}$




$\mathbf{w}^{(1)}$



$\mathbf{w}^{(2)}$




$\mathbf{w}^{(3)}$



⋮


$\mathbf{w}^{(c-1)}$



$\mathbf{w}^{(c)}$



Conv layer output




This can be used as the  
input to another layer



# Convolutional Neural Networks

## Activation layer


- It introduces non-linearity to the network
- Non-linearity helps the network learn complex patterns
- The most common activations are Sigmoid, ReLU and Tanh




# Convolutional Neural Networks

## Activation layer



$ReLU(x) = \max(0, x)$   
(Rectified Linear Unit)



$$\sigma(x) = \frac{1}{1 + e^{-x}}$$




$$\tanh(x)$$



# Convolutional Neural Networks

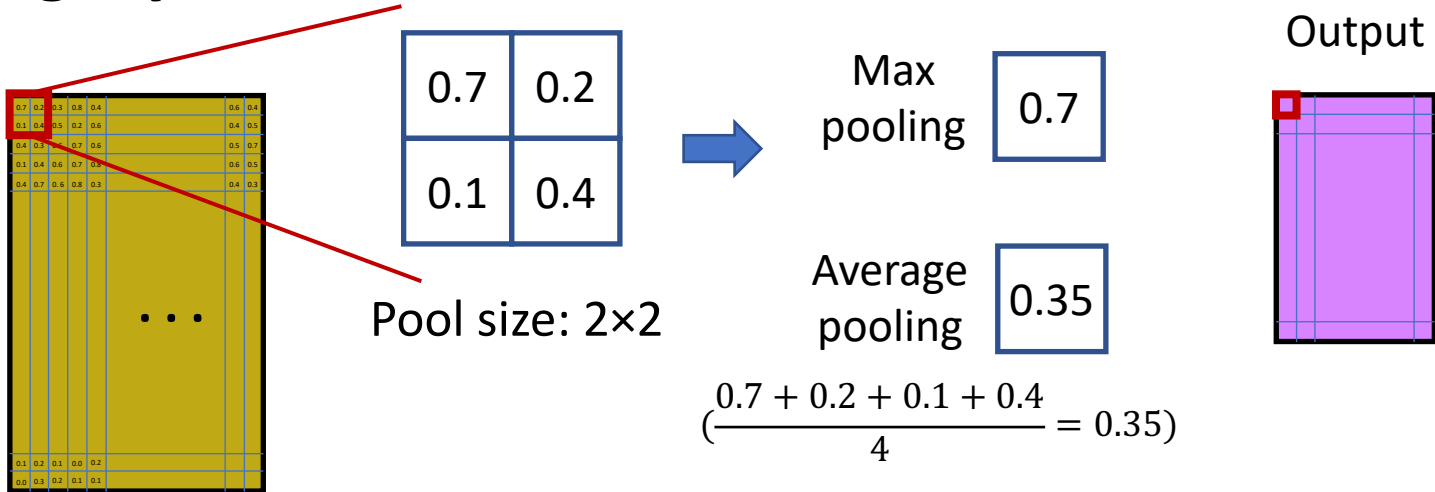
## Pooling layer

- It down-samples its input
- It reduces the network sensitivity to the features' location
- Two common pooling layers are average and max pooling



# Convolutional Neural Networks

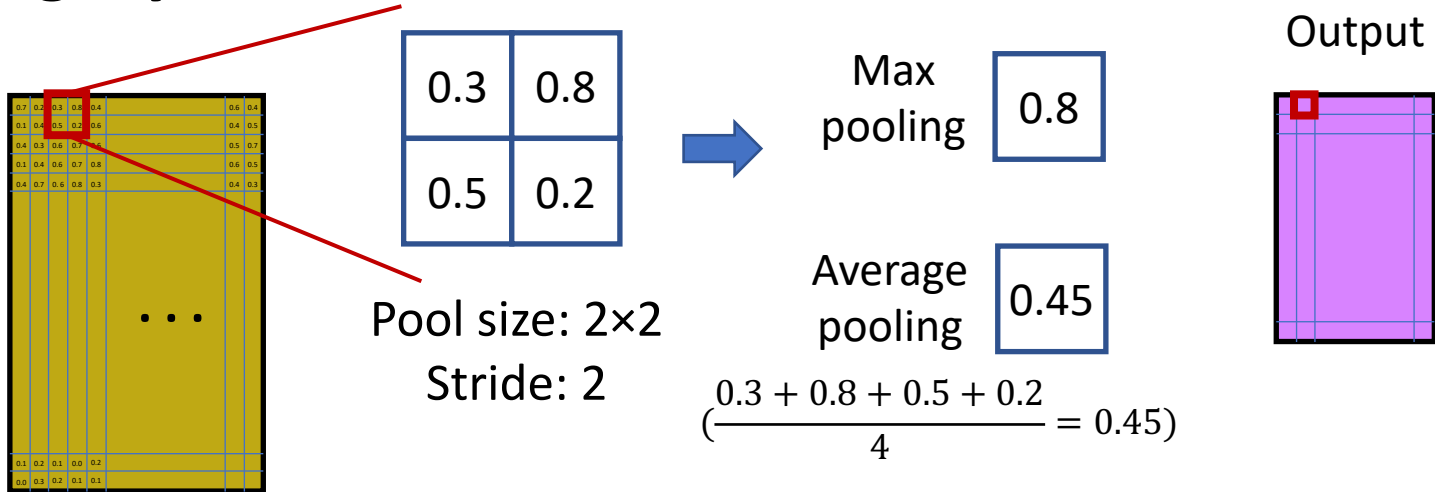
## Pooling layer






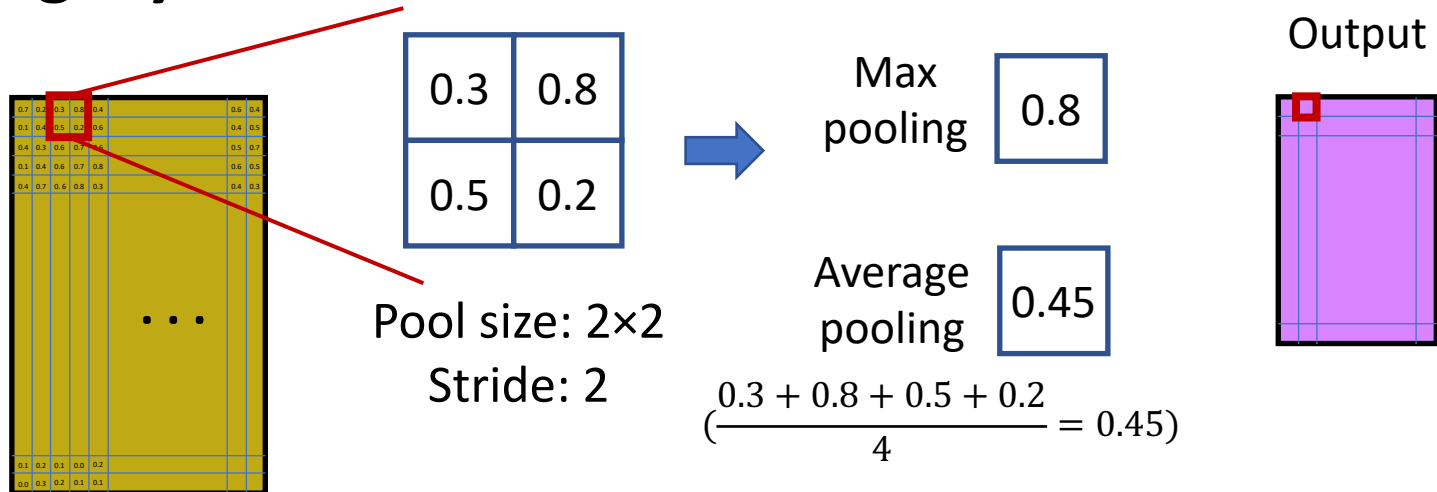
# Convolutional Neural Networks

## Pooling layer



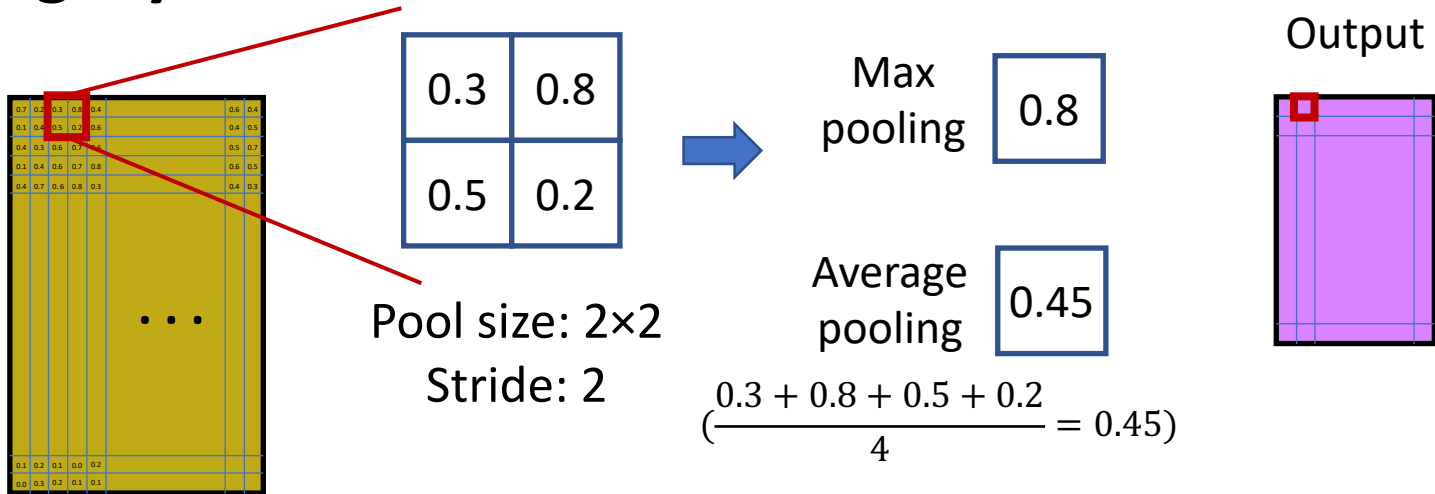
# Convolutional Neural Networks

## Pooling layer



# Convolutional Neural Networks

## Pooling layer




Pooling layer is not learnable

# Convolutional Neural Networks

## Fully-Connected layer


- It connects all its inputs to all the outputs
- It is usually used as the last layers of a network (before classification)



# Convolutional Neural Networks

## Fully-Connected layer


- It connects all its inputs to all the outputs
- It is usually used as the last layers of a network (before classification)



# Convolutional Neural Networks

## Drop-out layer


- It randomly drops out (sets to 0) neurons during training
- This prevents over-fitting (regularisation)



# Convolutional Neural Networks

## Batch-Normalisation layer

- It normalises its input across mini-batches
- It stabilises the network against the changes in the distribution of its input
- It helps the network generalise better to test data (regularisation)



# Convolutional Neural Networks

## Batch-Normalisation layer

It subtracts the mini-batch mean from its input and divides it by the mini-batch standard deviation

$$\hat{x} = \frac{x - m}{s}$$

This is then shifted and scaled by two learnable parameters

$$x^{BN} = \hat{x} \cdot \gamma + \beta$$

The BN output has a distribution with mean  $\beta$  and standard deviation  $\gamma$



# Popular CNN Architectures

# Popular CNN Architectures

## CNN – VGG16

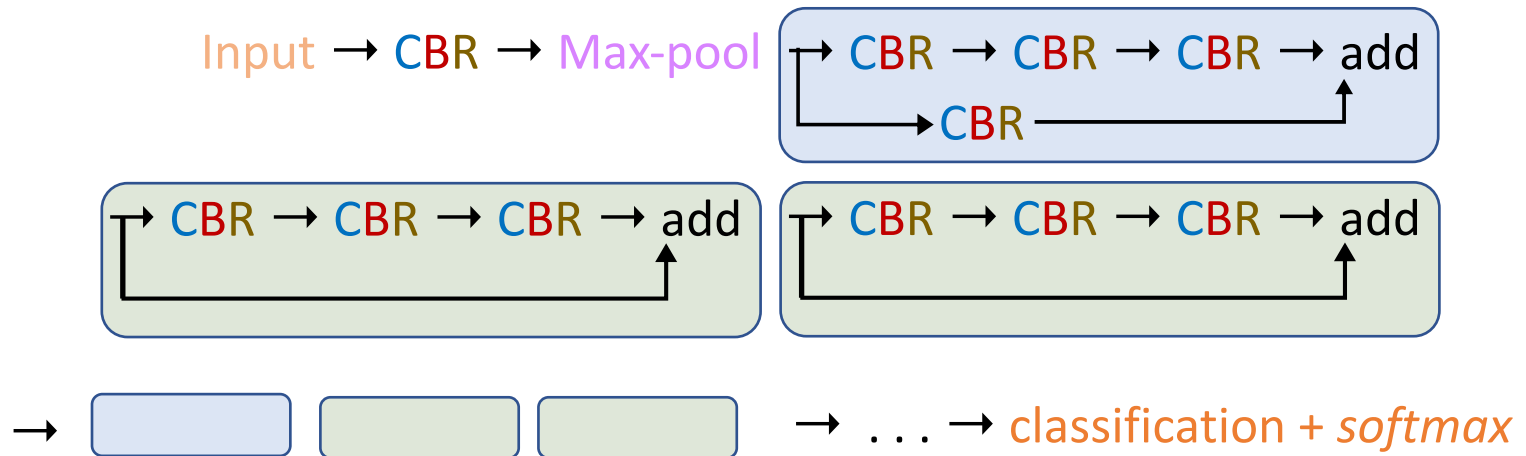
- It has 138 million parameters
- Its output layer has 1000 neurons

Input → Conv → ReLU → Conv → ReLU → Max-pool  
→ Conv → ReLU → Conv → ReLU → Max-pool  
→ Conv → ReLU → Conv → ReLU → Conv → ReLU → Max-pool  
→ Conv → ReLU → Conv → ReLU → Conv → ReLU → Max-pool  
→ FC → ReLU → FC → ReLU → classification + softmax

# Popular CNN Architectures

## CNN – ResNet

- It uses residual (shortcut) connections
- It has increased the number of layers (152)
- It has 26 million parameters (CBR: Conv → BN → ReLU)



# Popular CNN Architectures

## CNN for image classification


- AlexNet
- Inception-v1/ v2/ v4
- Inception-ResNets
- Xception
- DenseNet
- ResNeXt

# Other Popular Neural Networks

# Other Popular Neural Networks

## Recurrent Neural Networks (RNN)


- They process sequences of data
- They back-propagate loss through time



# Other Popular Neural Networks

## Long Short-Term Memory Networks (LSTM)


- They are a special type of Recurrent Networks
- They can store information in their memory for long durations



# Other Popular Neural Networks

## Auto-Encoders

- They reconstruct their input
- They are used for dimensionality reduction, denoising, ...






# Other Popular Neural Networks

## Generative Adversarial Networks (GAN)

- They are generative models – e.g. when trained they can generate fake images out of noise
- They have two sub-network (discriminative and generative)
- The two sub-networks compete against each other (adversarial)



# Summary and Conclusion

- We saw different types of layers in a CNN

Convolutional, Activation, Pooling

Fully-Connected, Drop-out, Batch Normalisation, ...

- We saw some popular neural networks
  - CNN (VGG and ResNet)
  - RNN and LSTM
  - Auto-Encoder
  - GAN

Thanks for your attention